



Introduction to OpenMP

Introduction to OpenMP

Presenter: Daniel Trahan

Contact: Daniel.Trahan@colorado.edu

Website: www.colorado.edu/rc

Sign in: <http://tinyurl.com/curc-names>

Slides: <https://github.com/ResearchComputing/OpenMP-Spring-2020>

Overview

- What is OpenMP
- Compiling and Running applications with OpenMP
- Compiler Directives and Parallel Regions
- Memory Management
- Synchronization and Loop Scheduling

Why Parallel?

- Modern CPUs are capped in clock speed...
- Instead of increasing clock speed, add more transistors!
- Multiple processors on a single chip.
- Parallelism required to utilize full power!

What is OpenMP

- OpenMP is a parallelization programming library for C, C++, and Fortran.
 - Built into GNU and Intel Compilers
 - Single Processor/Shared Memory based parallelization strategy
- OpenMP code follows a strategy of parallel 'blocks' of code
 - Serial blocks of code run on one core and parallel blocks of code run on multiple cores.
 - Blocks are specified with OpenMP compiler directives.

Why OpenMP

- Very simple to implement parallelization to an existing block of code.
- No Explicit Communication among processes
- Fork/Join Parallelism
 - Process begins in serial executing on a single core.
 - When reaching a parallel region, processes are forked.
 - At the end of the parallel region, processes are rejoined.

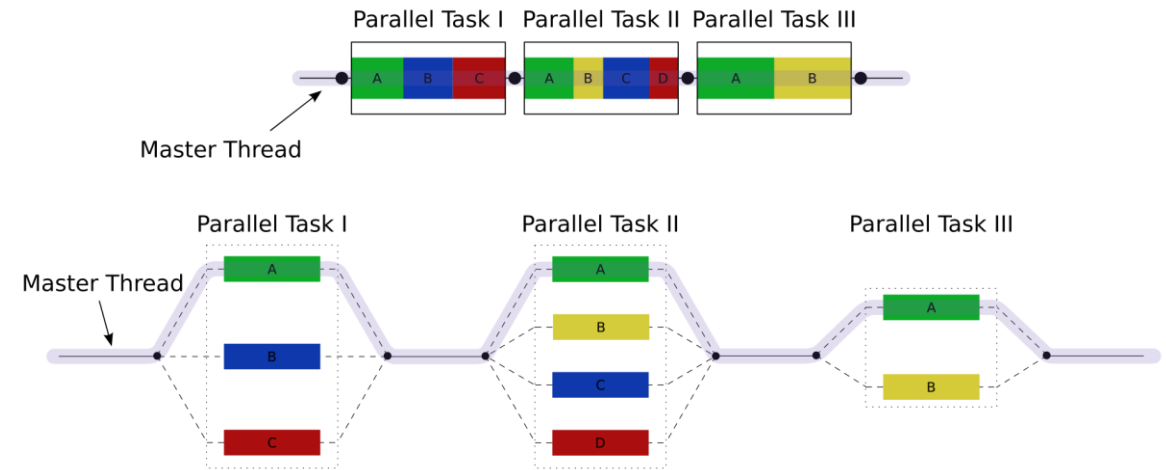


Image from: https://en.wikipedia.org/wiki/Fork%E2%80%93join_model

Getting started with OpenMP

- Declaring a Parallel Region:
 - Special comments in code are read by the compiler and parallelize the application.
 - 'parallel' directive
 - Any code within the code block is executed by all cores on your system.
 - Can be modified with environment variable.

```
export OMP_NUM_THREADS=<nthreads>
```

C/C++

```
#pragma omp parallel
{
    #parallelized code...
}
```

Fortran

```
use omp_lib
$omp parallel
    #parallelized code...
$omp end parallel
```

Compiling and Running with OpenMP

- Compiling a parallel application requires no special libraries or complex commands.
- Add a compiler flag!
 - With gcc: `-fopenmp`
 - With intel: `-qopenmp`

C/C++

```
gcc -fopenmp sample.c -o sample.exe  
#replace gcc w/ g++ for c++  
icc -qopenmp sample.c -o sample.exe  
#replace icc w/ icpc for c++
```

Fortran


```
gfortran -fopenmp sample.f90 -o sample.exe  
ifort -qopenmp sample.f90 -o sample.exe
```


Example 1: Parallel 'Hello World'

- Now that we know how to set up some basic parallelization, lets write a simple parallel hello world.
 - I will be using C for this example. The set up is pretty much the same for Fortran.
 - See [hello-omp.c](#) and [hello-omp.f90](#)

Compiler Directives (1)

- OpenMP understands a variety of other compiler directives with various clauses:
- Anatomy of Compiler Directives



- Ex: C/C++

```
#pragma omp for reduce(+:total)
```

Fortran

```
$!omp parallel shared(var1, var2)
```

Compiler Directives (2)

- A lot of different OpenMP directives!
- Can be broken into 3 different categories
 - Parallelization
 - Work-Sharing
 - Synchronization
- We will be talking about just 4 today:
 - 'Parallel', 'For', 'Section', and 'Barrier'

OpenMP Library Methods

- OpenMP includes several methods that can be utilized.
 - `omp_get_num_threads()` – Returns total number of threads
 - `omp_set_num_threads()` – Sets the number of available threads
 - `omp_get_thread_num()` – Returns the individual thread number
- Useful to manage parallelism and keep track of threads.
- Tons of more methods!

Directive: Parallel

- As shown previously, the OpenMP `parallel` directive is used to run a block of code by all available cores instead of a single core.
- Syntax:
- Must be invoked at the start of all parallel regions.
- Application will return to running in serial after completion.

C/C++

```
#pragma omp parallel  
{ # parallel code }
```

Fortran

```
!omp parallel  
$ parallel code  
!omp end parallel
```

Implicit Memory in OpenMP

- OpenMP has specific implicit rules that it utilizes to handle memory.
 - Any variables created outside the parallel scope are implicitly shared.
 - Any variables created inside the parallel scope are implicitly private.
 - Loop iteration variables are always private.
- Useful to know but can be confusing...

Explicit Memory in OpenMP (1)

- Memory in OpenMP are handled by compiler directive clauses.
 - ``shared`` variables have the same memory address on every thread.
 - ``private`` variables have different memory addresses for every thread.

C/C++

```
#pragma omp parallel shared(pub1, pub2, pub3) private(prv1, prv2, prv3)
{ # parallel code }
```

Fortran

```
$!omp parallel shared(pub1, pub2, pub3) private(prv1, prv2, prv3)
$ parallel code
$!omp end parallel
```

Explicit Memory in OpenMP (2)

- Can force explicit programming on all variables declared out of a scope.
 - ``default`` clause allows users to change behavior of all variables outside of a scope.
 - Accepts ``private``, ``shared``, or ``none`` as parameters.
 - ``none`` requires users manually classify all variables as shared or private.

C/C++

```
#pragma omp parallel default(none) shared(...) private(...)
```

Fortran

```
$!omp parallel default(none) shared(...) private(...)
```

Directive: For/Do

- The OpenMP ``for`/`do`` directive is a work sharing directive that divides a for loop among various workers.
- Syntax:
- Directive will split the next loop in the application among available workers
- Process is not threadsafe by default...

C/C++

```
#pragma omp for  
# for Loop
```

Fortran

```
$!omp do  
$ do Loop  
$!omp end do
```

Example 2:

- Let's parallelize a simple summation loop from all the numbers between 1 and 10000
- Prints out the total sum of the loop.
- Did our application properly compute the sum in parallel?
- Why or Why not?

Thread Safety

- Thread Safety is the assurance that a thread will execute without any unintentional side effects.
- Last example was overwriting a shared sum in a loop.
- Since the 'for' directive isn't thread safe, how can we ensure data integrity of our summation?

Reduction Clause

- With for/do loops that are all computing to a single reducible sum, we can use the ``reduction`` clause to avoid thread unsafe operations.
- Syntax:
- Requires a reduction operator and a reduction variable.
- Reduction variable is locally store and combined at the end.

C/C++

```
#pragma omp for reduction(+:reduce-var)  
# for Loop
```

Fortran

```
$!omp do reduction(+:reduce-var)  
$ do Loop  
$!omp end do
```


Directive: Barrier

- The OpenMP **barrier** directive is a synchronization directive that halts execution of threads until all threads reach the barrier line of code.
- Syntax:
- Useful for sections code that must be synchronized before execution.

C/C++

```
#pragma omp barrier
```

Fortran

```
$!omp barrier
```

Directive: Sections and Section (1)

- The OpenMP 'Sections' and 'Section' directives can be used to separate threads to individual tasks.
- The 'Sections' directive designates the body of directives that will be divided into individual chunks run by threads.
- The 'Section' directive designates actual code to be run by each thread.
- Each section is run by a single thread regardless of how many workers are available.

Directive: Sections and Section (2)

- Syntax:

C/C++

```
#pragma omp parallel sections
{
    #pragma omp section {
        # Thread 1 code
    }
    #pragma omp section {
        # Thread 2 code
    }
}
```

Fortran

```
$!omp parallel sections
$!omp section
$ Thread 1 code
$!omp section
$ Thread 2 code
$!omp end parallel sections
```

Example 3:

- Use the ``sections`` and ``section`` directive to distribute 3 loops to 3 threads that calculate the sum from 1 to 10, 1 to 100, and 1 to 1000.
- Use the reduction clause to reduce all 3 sums to one final sum and print out the result.

Additional Resources

- OpenMP Website: <https://www.openmp.org/>
- Jaka's Corner (OpenMP Ex): <http://jakascorner.com/blog/>
- Laurence Livermore: <https://computing.llnl.gov/tutorials/openMP/>
- GNU reference: <https://gcc.gnu.org/wiki/openmp>
- Intel reference: <https://software.intel.com/en-us/cpp-compiler-developer-guide-and-reference-openmp-support>

Questions and Thank You!

Contact: Daniel.Trahan@colorado.edu

Website: www.colorado.edu/rc

Sign in: <http://tinyurl.com/curc-names>

Slides: <https://github.com/ResearchComputing/OpenMP-Spring-2020>

Survey: <http://tinyurl.com/curc-survey18>